

Machine Learning (CE 40717)

Fall 2024

Ali Sharifi-Zarchi

CE Department
Sharif University of Technology

October 24, 2024



- 1 Optimization
- 2 The Loss Surface
- 3 Gradient Descent
- 4 Momentum
- 5 References

1 Optimization

Problem Definition

Convexity and Optimization

2 The Loss Surface

3 Gradient Descent

4 Momentum

5 References

1 Optimization

Problem Definition

Convexity and Optimization

2 The Loss Surface

3 Gradient Descent

4 Momentum

5 References

Problem Definition

- **Goal:** Given a function $f(x)$ of some variable x , find the value of x where $f(x)$ is **minimum** or **maximum**
- In neural networks, the goal is to make the prediction error as small as possible
- We want to find the network weights W^* that result in the **lowest loss**:

$$w^* = \arg \min_w J(w)$$

Problem Definition

- Simply put, we want to find the direction to step in that will **reduce our loss** as quickly as possible
- In other words, **which way is downhill?**



Figure 1: Visualization from CS231n, Stanford University

1 Optimization

Problem Definition

Convexity and Optimization

2 The Loss Surface


3 Gradient Descent

4 Momentum

5 References

Convexity

- **Definition:** A function is **convex** if, for any two points on the curve, the line segment connecting them lies above or on the curve
 - **Example:** Bowl-shaped curves
- **Convex functions** are easier to optimize because they have only **one global minimum** (the lowest point)
 - Analytical solutions ($\nabla f(x) = 0$) and second-order methods ($\nabla^2 f(x) > 0$) can be used for faster and more accurate convergence
 - **Gradient descent** is guaranteed to converge to the global minimum in convex functions



Convexity

- **Definition:** a function is **non-convex** if it has multiple local minima and maxima
 - **Global Minimum:** The very lowest point across the whole curve
 - **Local Minimum:** A point that's lower than nearby points, but not the lowest overall
 - **Saddle Points:** A flat region where the slope is almost zero. It can go up in some directions and down in others
- Non-convex functions make finding the **global minimum** complicated




Figure 2: Convex (above) and non-convex (below) functions. From 11-785 Introduction to Deep Learning, CMU, Fall 2024

- 1 Optimization
- 2 The Loss Surface
- 3 Gradient Descent
- 4 Momentum
- 5 References

Definition

- The loss surface represents how the error changes based on the network's weights
- The loss surface of a neural network is typically **non-convex** due to multiple layers, nonlinear activation functions, and complex interactions between parameters, resulting in multiple local minima and saddle points
- In large networks, most local minima give similar error values and are close to the **global minimum**. This isn't true for smaller networks




Figure 3: ResNet56 from <https://github.com/tomgoldstein/loss-landscape>

Loss Optimization

- How can we optimize a non-convex loss function?
- **Strategy:** Instead of randomly searching for a good direction, we calculate the **best direction** to reduce the loss
 - Mathematically, this is guaranteed to be the direction of the steepest descent




Figure 4: Gradient Descent visualization from <https://blog.skz.dev/gradient-descent>

1 Optimization

2 The Loss Surface

3 Gradient Descent

Recap of Gradient Descent

Problems with Gradient Descent

4 Momentum

5 References

Gradient Descent

- As introduced earlier, Gradient Descent is an iterative method for minimizing the error by updating network weights in the direction of the steepest descent (negative gradient)

$$w_{t+1} = w_t - \eta \nabla J(w_t)$$

where:

- η is called the 'learning rate' or 'step size'
- The goal is to keep adjusting the weights until we reach a point where the error is as low as possible (a local or global minimum)

Types of Gradient Descent

- **Batch Gradient Descent:** Uses the entire dataset to calculate the gradient. This gives smooth updates but can be slow
- **Stochastic Gradient Descent (SGD):** Uses one data point at a time, leading to faster but noisier updates
- **Mini-batch Gradient Descent:** Uses small groups (batches) of data points. This is a balance between batch and stochastic, combining speed with more stable updates

Types of Gradient Descent

Type	Advantages	Disadvantages
Batch	Stable convergence Accurate gradient estimate	Computationally expensive Slow for large datasets
Stochastic (SGD)	Fast updates Can escape local minima	Noisy updates May not converge smoothly
Mini-Batch	Faster than batch gradient descent, more stable than stochastic gradient descent (SGD) Efficient for larger datasets	Requires tuning batch size Some noise remains

Table 1: Comparison of Gradient Descent Types

- ## 5 References

Problem 1

- **SGD** is fast and can escape local minima, but it faces some issues
- What if the loss changes quickly in one direction but slowly in another?
 - **Slow** progress along the shallow dimension, **jitter** along the steep direction




Figure 5: SGD visualization from CS231n, Stanford University

Problem 2

- What if the loss function has a local minima or saddle point?
 - The algorithm may settle for **sub-optimal solutions** or take a **long time** to make significant progress




Figure 6: Saddle point from
https://en.wikipedia.org/wiki/Saddle_point

Problem 3

- Gradients that come from single data points or mini-batches can be **noisy**

Stochastic Gradient Descent




Figure 7: SGD visualization from <https://laptrinhx.com/understanding-optimization-algorithms-3818430905/>

Problem 4


- Why not just use mini-batches?
 - Even though mini-batch gradient descent helps **reduce the noise**, It can still converge **slowly** and might **get stuck** in regions where the gradients are small (like plateaus or valleys), which makes learning inefficient
- In addition, using the same learning rate for all dimensions can lead to:
 - Smooth convergence in some directions
 - Oscillations or divergence in other directions

Problem Definition

- So, how can we improve the vanilla Gradient Descent algorithm?
- **Proposal:**
 - Track oscillations in each direction
 - Increase steps in stable directions
 - Decrease steps in oscillating directions

Problem Definition

- **Goal:** Choose an appropriate learning rate to avoid slow convergence and getting stuck in local minima while not overshooting or becoming unstable
- **Naive Approach:** Test many different learning rates to find the one that works "just right"
 - This can be inefficient and time-consuming
- **Smarter Approach:** Design an adaptive learning rate that adapts to the loss surface
 - We can achieve this by incorporating **Momentum**



① Optimization

② The Loss Surface

③ Gradient Descent

④ Momentum

Momentum Definition and Types

Adam

⑤ References

1 Optimization

2 The Loss Surface

3 Gradient Descent

4 Momentum

Momentum Definition and Types


Adam

5 References

Momentum

- **Goal:** Speed up convergence by using past gradients to smooth out oscillations and avoid getting stuck
 - Continue moving in the same general direction as the previous steps
- **Benefit:** It accelerates learning in significant directions while reducing oscillations in less important ones

SGD without momentum



SGD with momentum




Figure 8: SGD comparison from <https://paperswithcode.com/method/sgd-with-momentum>

First Momentum

- The first momentum, denoted as m_t , is essentially **the moving average of the gradients**. The update rule is:

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

$$w_{t+1} = w_t - \eta m_{t+1}$$

where:

- m_{t+1} is the first moment at time step $t + 1$
- β_1 is the decay rate, controlling how much of the past gradients to include (typically 0.9 or 0.99)
- $\nabla_w J(w_t)$ is the gradient at time step t

Why use the first moment?

- Inspired by physics, it maintains movement due to accumulated momentum, similar to a ball rolling down a frictionless bowl

Second Momentum

- The second momentum, denoted as v_t , is a **moving average of the squared gradients**.
- It helps track the magnitude of the gradients over time. The update rule is:

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2)(\nabla_w J(w_t))^2$$

$$w_{t+1} = w_t - \eta v_{t+1}$$

where:

- v_{t+1} is the second moment at time step $t + 1$
- β_2 is the decay rate
- $\nabla_w J(w_t)$ is the gradient at time step t

Why use the second moment?

- The second moment helps **regulate update sizes** by adjusting for consistently large or small gradients, preventing overshooting or slow learning

Moment Bias Correction

- **Problem:** When we start training, both m_t and v_t are initialized to zero, causing their estimates to be biased toward zero in the early steps, especially when gradients are small.
- **Solution:** We use bias-corrected versions of m_t and v_t to address this:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- These corrections compensate for the bias by scaling m_t and v_t upward, especially in the early steps when t is small, ensuring more accurate estimates of the moments.

① Optimization

② The Loss Surface

③ Gradient Descent

④ Momentum

Momentum Definition and Types

Adam

⑤ References

Adam

- **Adaptive Moment Estimation (Adam)** combines **momentum** and **adaptive learning rates** by maintaining an exponentially decaying average of both past gradients and squared gradients
- Adam adjusts the learning rate for each parameter based on the gradient history
 - Larger gradients result in smaller update steps, and vice versa

Adam

- The update rule for the Adam optimizer at step $t + 1$ is denoted by w_{t+1} :

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2) (\nabla_w J(w_t))^2$$

$$\hat{m}_{t+1} = \frac{m_{t+1}}{1 - \beta_1^{t+1}}, \quad \hat{v}_{t+1} = \frac{v_{t+1}}{1 - \beta_2^{t+1}}$$

$$w_{t+1} = w_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

where:

- m_{t+1} and v_{t+1} are the first and second moments at step $t + 1$
- β_1 and β_2 are the decay rates for the first and second moments
- \hat{m}_{t+1} and \hat{v}_{t+1} are bias-corrected estimates of the first and second moments
- ϵ is a small constant to prevent division by zero

Comparison of Momentum Methods




Figure 9: GD comparison from <https://github.com/ilguyi/optimizers.numpy>




Figure 10: GD comparison on MNIST from kingma2014adam

- 1 Optimization
- 2 The Loss Surface
- 3 Gradient Descent
- 4 Momentum
- 5 References