

Machine Learning (CE 40717)

Fall 2024

Ali Sharifi-Zarchi

CE Department
Sharif University of Technology

October 26, 2024



- 1 Optimization
- 2 The Loss Surface
- 3 Newton's optimization Method
- 4 Gradient Descent
- 5 Momentum

1 Optimization

2 The Loss Surface

3 Newton's optimization Method

4 Gradient Descent

5 Momentum

Optimization Problem

- **Goal:** Find the value of x where $f(x)$ is at a **minimum** or **maximum**.
- In neural networks, we aim to minimize **prediction error** by finding the optimal weights w^* :

$$w^* = \arg \min_w J(w)$$

- Simply put: determine the **direction to step** that will quickly **reduce loss**.

Convexity and Optimization

- **Convex Functions:**

- A function is **convex** if any line segment between points on the curve lies **above or on** the curve.
- Convex functions are easier to optimize, as they have a single **global minimum**.
- **Gradient Descent** is guaranteed to reach the global minimum in convex functions.




Figure 1: Example of convex function (bowl shape)

Non-Convex Functions and Challenges

- **Non-Convex Functions:**
 - Characterized by multiple **local minima** and **saddle points**.
 - **Global Minimum:** Overall lowest point.
 - **Local Minimum:** Lower than nearby points, but not the lowest overall.
 - **Saddle Points:** Regions where the gradient is close to zero but can increase or decrease in other directions.
- Finding the **global minimum** is more complex in non-convex functions.




Figure 2: Convex (top) vs. Non-Convex (bottom) functions (CMU, 11-785)

- A set of navigation icons typically found in Beamer presentations, including symbols for back, forward, search, and other slide controls.

Loss Surface Definition

- The **loss surface** shows how error changes based on network weights.
- For neural networks, the loss surface is typically **non-convex** due to multiple layers, nonlinear activations, and complex parameter interactions, resulting in **multiple local minima** and **saddle points**.
- In large networks, most local minima yield similar error values close to the **global minimum**; this is less true in smaller networks.




Figure 3: Loss surface of ResNet56. Source: <https://github.com/tomgoldstein/loss-landscape>

Loss Optimization

- **Goal:** How can we optimize a non-convex loss function effectively?
- **Newton's Method:**
 - When optimizing, we look for critical points where the derivative $f'(x) = 0$, which may indicate minima, maxima, or saddle points.
 - Newton's Method uses the second derivative (Hessian) to adjust step sizes, which can lead to faster convergence compared to Gradient Descent.
- **Gradient Descent:**
 - This method identifies the steepest descent direction to guide the optimization process.

- 1 Optimization
- 2 The Loss Surface
- 3 Newton's optimization Method**
Newton's Method
- 4 Gradient Descent
- 5 Momentum

Newton's Method

- Newton method is originally intended to **find the root(s)** of an equation.
- **Example:** for the equation $x^2 - 1 = 0$, we can find the roots by decomposing $(x - 1)(x + 1) = 0$ which gives $x = 1, x = -1$
- **But, what about complex equations?**
 - We can use **numerical method** to find the root of an equation, one of them is by using **Newton's method**

Definition

- **Objective:** Derive Newton's method by finding the tangent line of $f(x)$ at x_0 .
- **Tangent Line Equation:** Given a point x_0 where $f(x_0) \neq 0$, the tangent line at x_0 is:

$$y = mx_0 + c$$

- **Gradient:** The slope m matches the derivative of $f(x)$ at x_0 :

$$m = f'(x_0)$$




Figure 4: Finding root location by using Newton's method. Source: <https://ardianumam.wordpress.com>

Formulating the Tangent Line

- **Finding c :** Substitute $(x_0, f(x_0))$ into $y = mx + c$, where $y = f(x_0)$ and $m = f'(x_0)$:

$$f(x_0) = f'(x_0)x_0 + c \Rightarrow c = f(x_0) - f'(x_0)x_0$$

- **Tangent Line Equation:** Substitute $m = f'(x_0)$ and c back:

$$y = f'(x_0)x + f(x_0) - f'(x_0)x_0$$

- Simplify to get:

$$y = f(x_0) + f'(x_0)(x - x_0)$$

Newton's Iterative Step

- To approximate the root, set $y = 0$ in the tangent equation:

$$0 = f(x_0) + f'(x_0)(x_1 - x_0)$$

- Rearrange to solve for x_1 :

$$x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}$$

- **Iteration:** Repeat this step to approximate the root.

Newton's Method for Optimization

- Newton's method for finding roots is based on a first-order approximation (tangent line).
- For optimization, we use a second-order Taylor approximation to find the minimum.
- **Second-order Taylor expansion** of $f(x)$ around $x = x_0$:

$$f(x) \approx f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2}(x - x_0)^2$$

- Rearranged for minimal value location:

$$f(x) \approx \frac{1}{2}f''(x_0)x^2 + [f'(x_0) - f''(x_0)x_0]x + [f(x_0) - f'(x_0)x_0 + \frac{1}{2}f''(x_0)x_0^2]$$

Deriving the Update Formula for Minimization

- To locate the minimum, take the derivative with respect to x and set it to zero:

$$\frac{d}{dx}f(x) \approx f''(x_0)x + [f'(x_0) - f''(x_0)x_0] = 0$$

- Solving for x yields:

$$x = x_0 - \frac{f'(x_0)}{f''(x_0)}$$

- This is the update step for Newton's method in optimization, guiding us to the minimum. The general update rule is:

$$w_{t+1} = w_t - H^{-1} \nabla_w J(w_t)$$




Figure 5: Finding root using Taylor's expansion and Newton's method

Newton's Method: Advantages and Disadvantages

- Newton's method offers various benefits but also has limitations, especially in large-scale machine learning. Below is a summary:

Advantages	Disadvantages
Faster Convergence Quadratic convergence enables reaching minima faster in convex problems.	Computationally Expensive Requires Hessian calculation, making it costly in high-dimensional models.
Adaptive Step Sizes Curvature-based step adjustment avoids slow progress in shallow regions.	Memory Intensive Storing the Hessian matrix is memory-intensive for models with millions of parameters.
Reduced Oscillations Curvature information stabilizes paths in oscillatory regions.	Convergence Challenges May converge to saddle points in non-convex functions common in machine learning.

- 1 Optimization
- 2 The Loss Surface
- 3 Newton's optimization Method
- 4 Gradient Descent
- 5 Momentum

Gradient Descent Overview

- **Gradient Descent:** As mentioned earlier in this course, Gradient Descent is an iterative method to minimize error by updating weights in the direction of the **negative gradient**:

$$w_{t+1} = w_t - \eta \nabla J(w_t)$$

where η is the **learning rate**.

- **Types of Gradient Descent:**
 - **Batch:** Full dataset for stable but slow updates.
 - **Stochastic (SGD):** One data point for fast, noisy updates.
 - **Mini-Batch:** Small batches, balancing speed and stability.

Problems with Gradient Descent

- **High Variability (SGD):** Quick in steep directions but slow in shallow ones, causing jitter and slow progress.
- **Local Minima and Saddle Points:** Risk of **sub-optimal solutions** or long convergence times in flat regions.
- **Noisy Updates:** Using individual points or mini-batches introduces noise, affecting stable convergence.




Figure 6: SGD Variability (CS231n, Stanford)




Figure 7: Saddle Point (en.wikipedia.org)

1 Optimization

2 The Loss Surface

3 Newton's optimization Method

4 Gradient Descent

5 Momentum

First Moment (Momentum)

Second Moment (Variance)

Adam: Adaptive Moment Estimation

Problem Definition and Solution

- **Objective:** Enhance the vanilla Gradient Descent algorithm to improve convergence and stability.
- **Challenges:**
 - Selecting an appropriate learning rate is crucial to avoid slow convergence and getting stuck in local minima.
- **Proposed Solution:**
 - Instead of testing multiple learning rates, which can be inefficient and time-consuming, incorporate **Momentum** to adaptively adjust the learning rate based on oscillations:
 - Increase steps in stable directions.
 - Decrease steps in oscillating directions.




Figure 8: Momentum smooths oscillations and accelerates progress. Source: paperswithcode

Introduction to Momentum in Optimization

- **Origin of Momentum:**

- Inspired by Newtonian physics, momentum in optimization uses the concept of velocity in motion, accumulating gradient history to smooth the learning trajectory, akin to an object moving based on past inertia.
- Initially introduced to tackle challenges in gradient descent, where inconsistent gradients or noisy updates lead to erratic and slow convergence.

- **Purpose of Momentum:**

- **Dampens Oscillations:** Utilizes prior gradients to minimize oscillations along steep or erratic regions, resulting in a smoother and more stable path.
- **Speeds Up Convergence:** Particularly effective in narrow valleys or flat regions, where standard gradient descent may struggle or oscillate, causing slow progress.

1 Optimization

2 The Loss Surface

3 Newton's optimization Method

4 Gradient Descent

5 Momentum

First Moment (Momentum)

Second Moment (Variance)

Adam: Adaptive Moment Estimation

First Moment (Momentum)

- **Definition:** The first moment, m_t , represents a moving average of past gradients. It builds "velocity" that propels learning in a consistent direction.
- **Update Rule:**

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

$$w_{t+1} = w_t - \eta m_{t+1}$$

where:

- β_1 : Decay rate, usually 0.9 or 0.99, which controls the weight of past gradients.
- η : Learning rate.
- **Why Use First Momentum?**
 - Inspired by the idea of rolling momentum, it smooths and accelerates learning by sustaining direction from prior gradients.
 - This type of momentum is ideal for traversing narrow valleys or regions where standard gradient descent would oscillate.

Example of First Moment




Figure 9: Stochastic gradient descent without momentum stops at a local minimum. Source: Akash Ajagekar (SYSEN 6800 Fall 2021)




Figure 10: Stochastic gradient descent with momentum stops at the global minimum. Source: Akash Ajagekar (SYSEN 6800 Fall 2021)

1 Optimization

2 The Loss Surface

3 Newton's optimization Method

4 Gradient Descent

5 Momentum

First Moment (Momentum)

Second Moment (Variance)

Adam: Adaptive Moment Estimation

Second Moment (Variance)

- **Definition:** The second moment, v_t , represents the moving average of squared gradients. It measures the gradient magnitude over time.
- **Update Rule:**

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2)(\nabla_w J(w_t))^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{v_{t+1} + \epsilon}} \nabla_w J(w_t)$$

where:

- β_2 : Decay rate for variance (usually 0.99 or 0.999).
- ϵ : Small constant to prevent division by zero.
- **Why Use Second Momentum?**
 - Adjusts step size based on gradient magnitude, preventing large steps when gradients are large and accelerating learning when they are small.

Moment Bias Correction

- **Problem:** When we start training, both m_t and v_t are initialized to zero, causing their estimates to be biased toward zero in the early steps, especially when gradients are small.
- **Solution:** We use bias-corrected versions of m_t and v_t to address this:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- These corrections compensate for the bias by scaling m_t and v_t upward, especially in the early steps when t is small, ensuring more accurate estimates of the moments.

1 Optimization

2 The Loss Surface

3 Newton's optimization Method

4 Gradient Descent

5 Momentum

First Moment (Momentum)

Second Moment (Variance)

Adam: Adaptive Moment Estimation

Introduction to Adam Optimizer

- **Origin and Purpose:**
 - Proposed in 2014 by Diederik Kingma and Jimmy Ba, Adam (Adaptive Moment Estimation) addresses key limitations in earlier optimization methods by combining aspects of **momentum** and **adaptive learning rates**.
 - Adam is designed to handle sparse gradients and noisy updates by adjusting the learning rate for each parameter based on historical gradients.
- **Core Idea:**
 - Adam optimizes by maintaining two moving averages — the **first moment (mean of gradients)** and the **second moment (variance of gradients)** — allowing it to adapt learning rates for each parameter individually.

Adam's Adaptive Learning Rate Mechanism

- **Why Adaptive Rates?**
 - Unlike traditional SGD, Adam adapts the learning rate for each parameter based on recent gradient magnitudes.
 - Large gradients lead to reduced update sizes, while smaller gradients allow larger updates, balancing convergence speed and stability.
- **Moment Tracking**
 - The **first moment** (m_t) tracks the mean of gradients to provide momentum.
 - The **second moment** (v_t) tracks squared gradients, enabling Adam to normalize updates and prevent sudden changes in direction.

Mathematical Formulation of Adam

- **Adam Update Rules:**

- First moment estimate:

$$m_{t+1} = \beta_1 m_t + (1 - \beta_1) \nabla_w J(w_t)$$

- Second moment estimate:

$$v_{t+1} = \beta_2 v_t + (1 - \beta_2) (\nabla_w J(w_t))^2$$

- Bias-corrected moments to address initialization bias:

$$\hat{m}_{t+1} = \frac{m_{t+1}}{1 - \beta_1^{t+1}}, \quad \hat{v}_{t+1} = \frac{v_{t+1}}{1 - \beta_2^{t+1}}$$

- Update step for parameter w_t :

$$w_{t+1} = w_t - \eta \frac{\hat{m}_{t+1}}{\sqrt{\hat{v}_{t+1} + \epsilon}}$$

Adam Pseudo-code

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

Figure 11: Adam Pseudo-code. Source: kingma2014adam

Comparison of Momentum Methods




Figure 12: GD comparison from
<https://github.com/ilguyi/optimizers.numpy>




Figure 13: GD comparison on MNIST from
kingma2014adam