

Machine Learning (CE 40477)

Fall 2024

Ali Sharifi-Zarchi

CE Department
Sharif University of Technology

November 21, 2024



- 1 Channels
- 2 Pooling
- 3 Receptive field & inductive bias
- 4 References

4 References

Channels

Previously, we discussed 2D inputs; however, although images are inherently 2D, they need to be represented as **3D matrices** to display color information.

- Pixel values range from 0 to 255.
- It is not possible to represent all the colors in a picture using only a single channel of numbers from 0 to 255.
- Thus, we represent them using **three channels**.




Figure adapted from Source

Channels

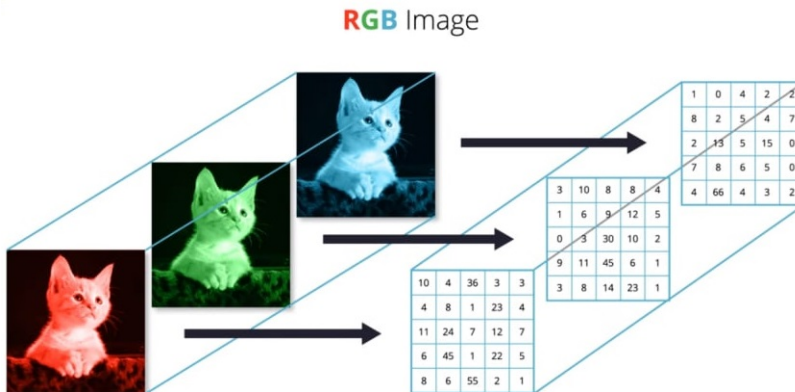


Figure adapted from Source

Channels

- Each filter produces **one output channel**. By applying multiple filters, we can create multiple output channels, allowing each channel to learn **distinct** features.




Figure adapted from Source

Let's take a closer look at the calculations

Channels

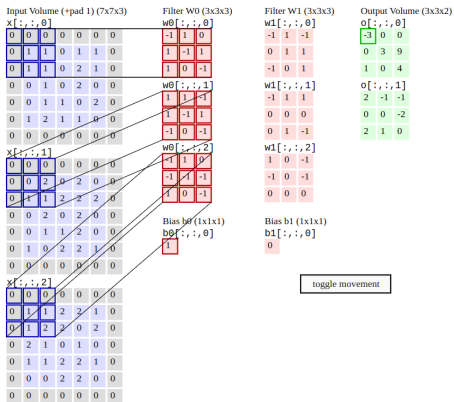


Figure adapted from [2]

Channels

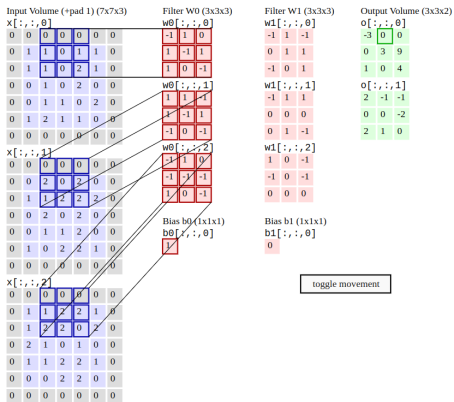


Figure adapted from [2]

Channels

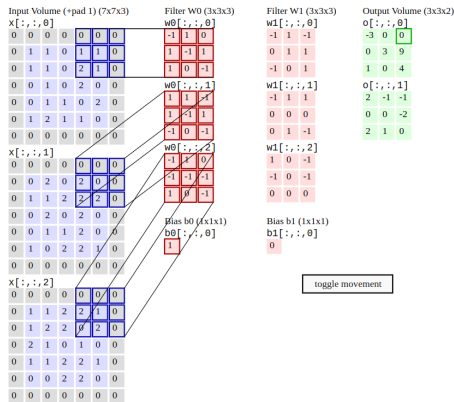


Figure adapted from [2]

Channels

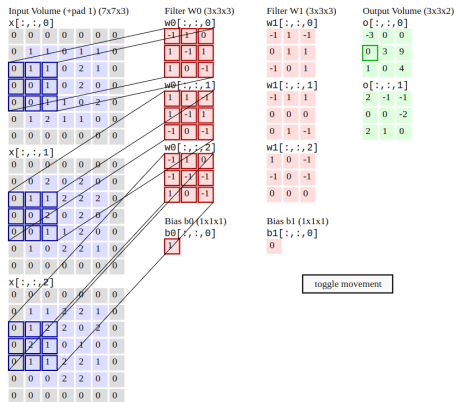


Figure adapted from [2]

Channels

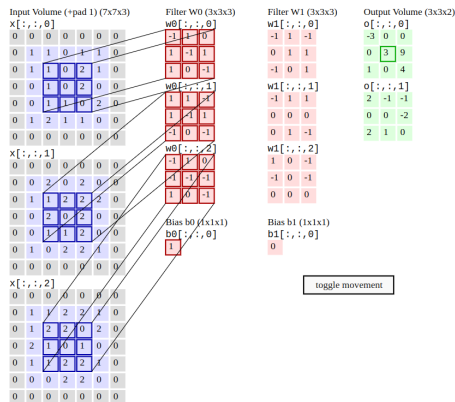


Figure adapted from [2]

Channels




Figure adapted from [2]

Channels




Figure adapted from [2]

Channels

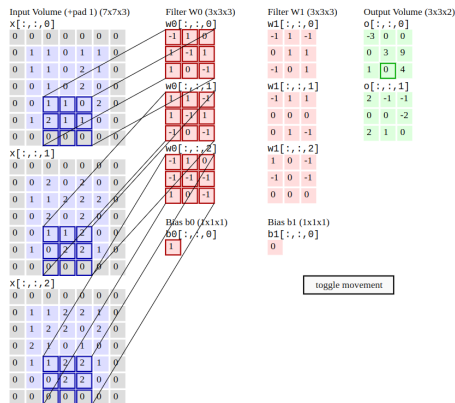


Figure adapted from [2]

Channels

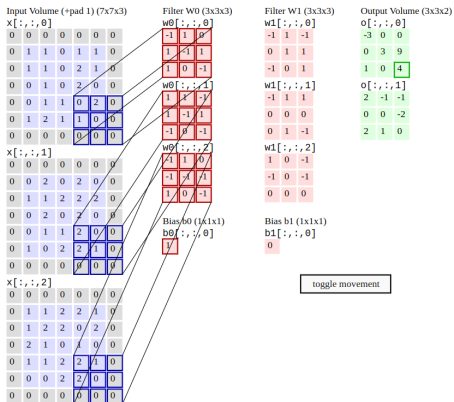


Figure adapted from [2]

Channels

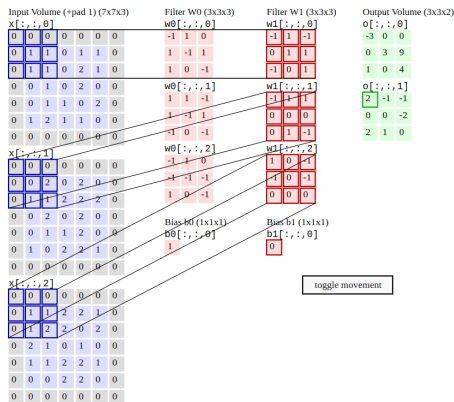


Figure adapted from [2]

Channels

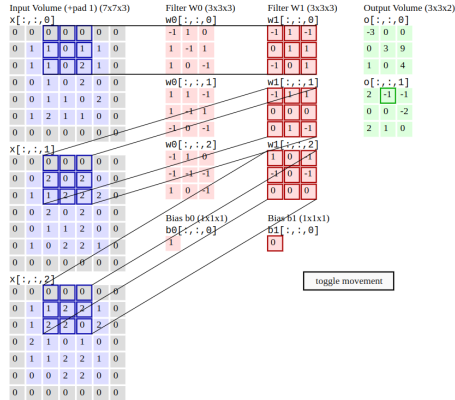


Figure adapted from [2]

Channels

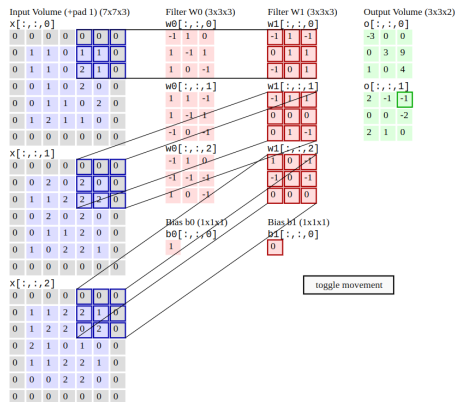


Figure adapted from [2]

Channels

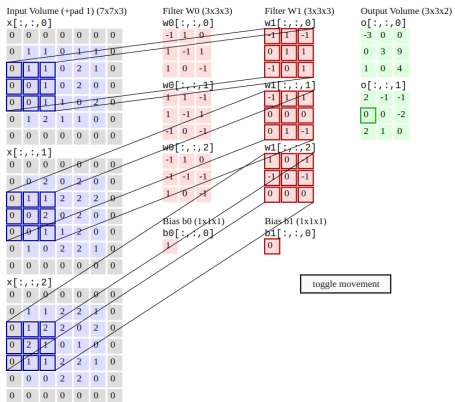


Figure adapted from [2]

Channels

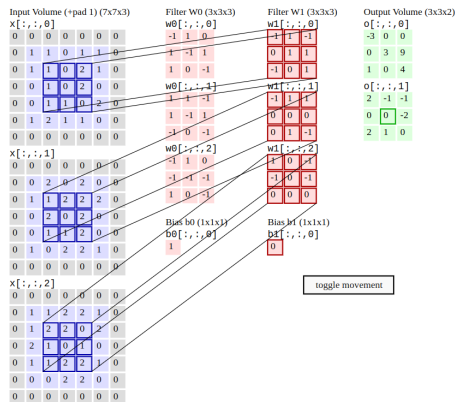


Figure adapted from [2]

Channels

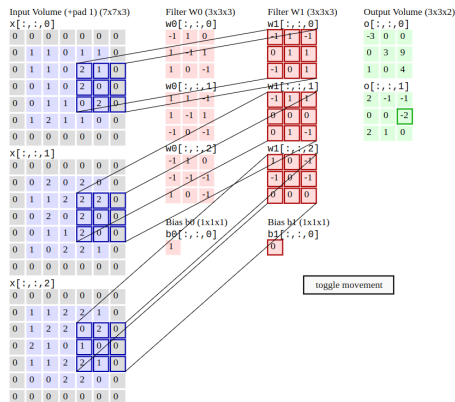


Figure adapted from [2]

Channels

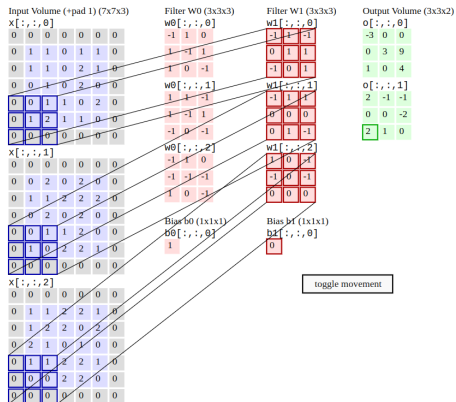


Figure adapted from [2]

Channels

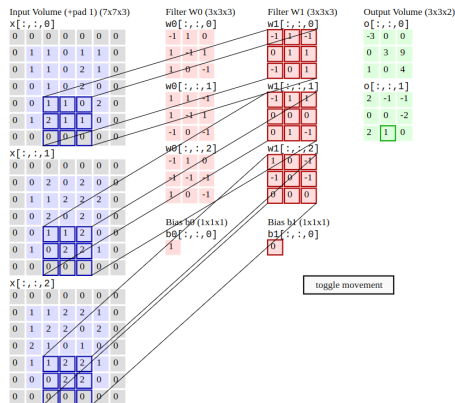


Figure adapted from [2]

Channels

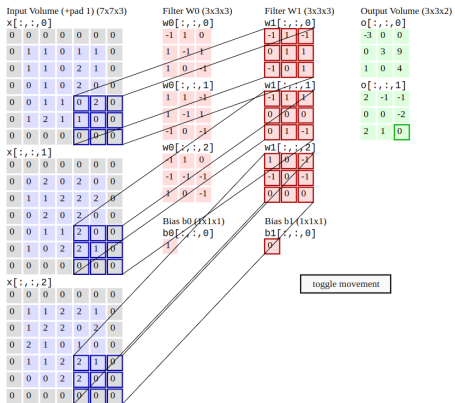
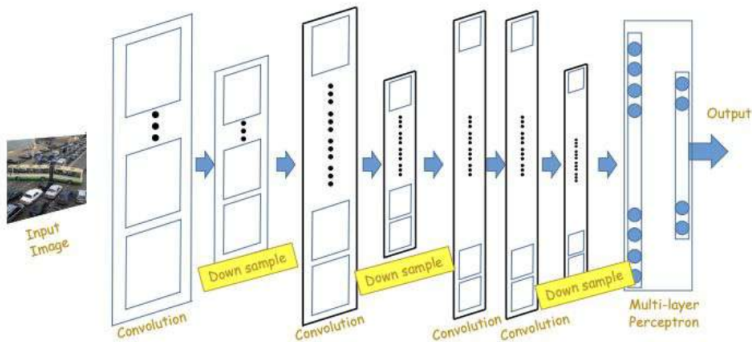


Figure adapted from [2]


- 1 Channels
- 2 Pooling
- 3 Receptive field & inductive bias
- 4 References

Pooling

- Convolution and activation layers are often followed by pooling layers intermittently.
 - Pooling layers often alternate with convolution layers, although this is not mandatory.



Pooling Type




Two Primary Types of Pooling:

- **Max Pooling:** Selects the **maximum** value from each section of the feature map.
- **Min Pooling:** Selects the **minimum** value from each section of the feature map.
- **Average Pooling:** Calculates the **average** value for each section of the feature map.


Figure adapted from source

Max Pooling



- This is the most common type of pooling layer.
- Provides invariance to small translations.

Downsampling




- Downsampling can be done by a simple convolution layer with stride larger than 1, Replacing the max pooling layer with a convolutional layer.

Question

- What if we want to increase the dimensions?

Up-sampling CNN

- Resizing feature maps is a **common** operation in neural networks, especially those used for image segmentation tasks.
- This architecture is often referred to as an **Encoder-Decoder** network.



Nearest Neighbors

- **Nearest Neighbors:** Nearest Neighbors involves copying an input pixel value to the K -nearest neighboring pixels, with K based on the expected output.



1	2
3	4

Input: 2 x 2




1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Output: 4 x 4

Bilinear Interpolation

- **Bilinear Interpolation:** In Bilinear Interpolation, the four nearest pixel values are used to compute a weighted average based on their distances, resulting in a smoothed output.



Bed Of Nails

- **Bed Of Nails:** In this method, the input pixel value is copied to the corresponding position in the output image, with zeros filling the remaining positions.

“Bed of Nails”

1	2
3	4

Input: 2 x 2




1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Output: 4 x 4

Max-Unpooling

- **Max-Unpooling:** In max-unpooling, the index of the **maximum value** is saved for each max-pooling layer during encoding. During decoding, the saved index is used to map the input pixel to its original position, with zeros filling all other positions.



Backpropagation With Pooling Layers

In our previous discussions, we explored the process of backpropagation in CNNs without considering pooling layers. Now, let's discuss how to adapt our algorithm to include pooling layers.

- The primary task is to handle the gradients effectively during backpropagation through pooling layers.
- In both cases, the gradients from the next layer are **passed back** to the previous layer through the pooling operation.


Case 1




Case 1

For **Max Pooling**:

- The gradient is propagated only through the **indices of the maximum values** identified during the forward pass.
- All other positions receive a gradient of **zero**.




Case 2



Case 2


For **Average Pooling**:

- The gradients are **uniformly distributed** across the pooled region.



- 1 Channels
- 2 Pooling
- 3 Receptive field & inductive bias
- 4 References

Receptive Field



- **Receptive Field:** How large is the region in the **input or previous layer** seen by a neuron on the n -th convolutional layer?
- In a convolution with kernel size K , each element in the next layer is based on a $K \times K$ **receptive field** from the previous layer.

Figure adapted from [3]

Power Of Small Filters

Assuming an input size of $H \times W \times C$, and convolutions are used with C filters to preserve depth (stride 1, with padding to maintain H and W dimensions).

one CONV with 7 x 7 filters

Number of weights

$$= C \times (7 \times 7 \times C) = 49C^2$$

three CONV with 3 x 3 filters

Number of weights

$$= 3 \times C \times (3 \times 3 \times C) = 27C^2$$

Both options achieve a receptive field of 7. However, using **multiple smaller filters** reduces the number of **parameters**, introduces more **nonlinearity**, and generally leads to a **more efficient**, expressive model.

Inductive Bias In CNNs

Inductive Bias:

The assumptions a model uses to generalize from training data to new, unseen data.

Key Features of Inductive Bias in CNNs:

- **Weight Sharing:**

A single filter is applied across various regions of the input, significantly decreasing the parameter count.

- **Locality:**

CNNs utilize small filters (e.g., 3×3) that concentrate on local regions, aligning well with image data where local structures are significant.

- CNNs are more sample-efficient than FCNs due to their inductive biases.

- 1 Channels
- 2 Pooling
- 3 Receptive field & inductive bias
- 4 References

Contributions

- **This slide was prepared with contributions from:**
 - Ali Aghayari
 - Behrooz Azarkhalili
 - Arian Amani
 - Hamidreza Yaghoubi

- [1] M. Soleymani Baghshah, “Deep learning.” Lecture slides.
- [2] F.-F. Li, Y. Li, and R. Gao, “Cs231n: Deep learning for computer vision.” Lecture slides.
- [3] B. Raj, R. Singh, and B. Dhingra, “11-785: Introduction to deep learning.” Lecture slides.
- [4] M. Kellis, “6.874 computational systems biology: Deep learning in the life sciences.” Lecture slides.
- [5] H. Li, “Eleg 5491: Introduction to deep learning.” Lecture slides.
- [6] M. Elgendy, *Deep Learning for Vision Systems*.
Manning Publications, 2020.
- [7] DeepMind, “Deep learning for ai with geoffrey hinton, yoshua bengio, and yann lecun.” YouTube video.