

Machine Learning (CE 40717)

Fall 2024

Ali Sharifi-Zarchi

CE Department
Sharif University of Technology

November 9, 2024

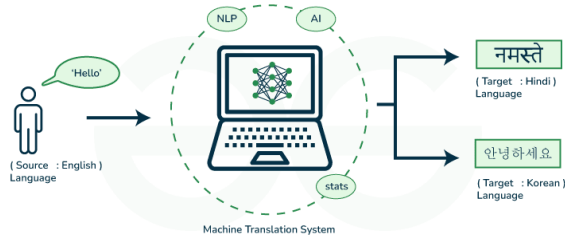


3 References

3 References

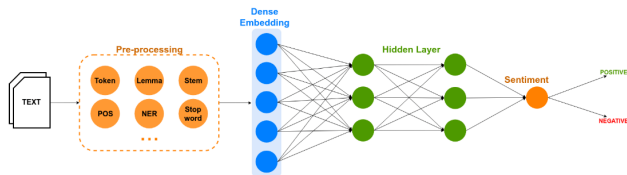
Natural Language Processing

- Language is central to human interaction; many of our daily activities revolve around text and language.
- Natural Language Processing (NLP) enables computers to understand and generate human language.



- NLP helps translate text from one language to another.

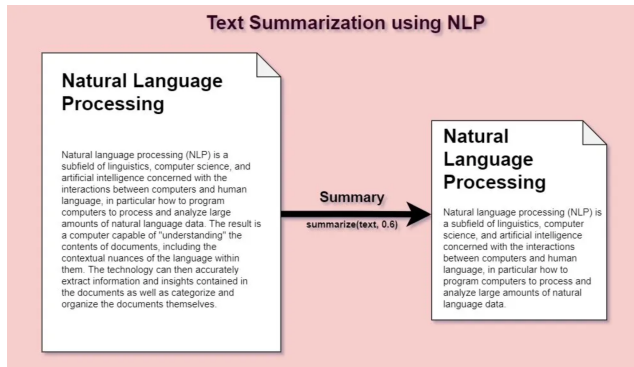
Figure adapted from [geeksforgeeks.org/machine-translation-of-languages-in-artificial-intelligence/](https://www.geeksforgeeks.org/machine-translation-of-languages-in-artificial-intelligence/)



- Determines the sentiment (e.g., positive or negative) expressed in a text.

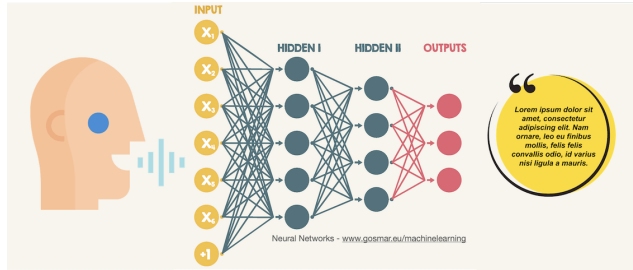
Figure adapted from www.mdpi.com/2079-9292/9/3/483

Text Summarization



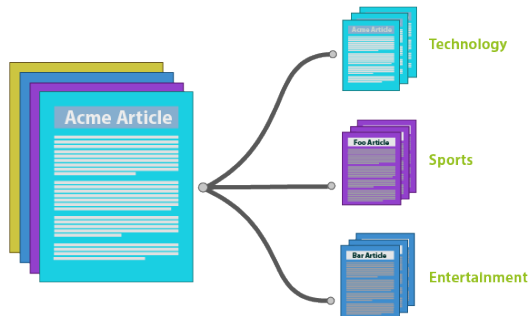
- Automatically generates a concise summary of longer text.

Speech Recognition



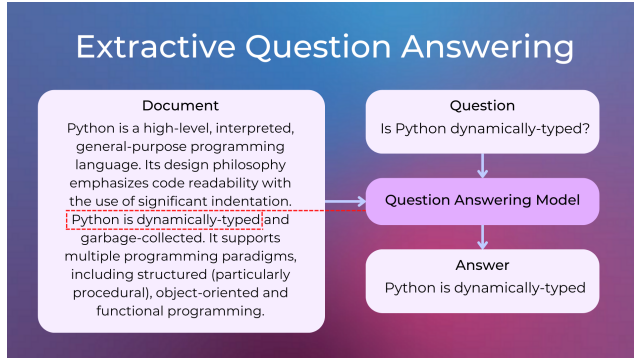
- Converts spoken language into text.

Text Classification



- Categorizes text into predefined groups or topics.

Extractive Question Answering



- Answers questions based on a given text or dataset.

- To process text effectively, the first step is to represent words in a way that models can understand.
- We need to transform words into vectors or dense representations to capture their meaning and relationships.
- This is crucial for enabling machines to understand and use language as humans do.

- Traditional models use methods like one-hot encoding, which lacks semantic understanding and cannot capture relationships between words.
- We need better word representations that are both dense and semantic.

Definition of One-Hot Encoding

- One-hot encoding is a straightforward method for representing categorical data, such as words, as discrete vectors.
- Each word is represented as a binary vector with the same length as the vocabulary size.
- All vector elements are set to 0 except for one position, which is set to 1, identifying the word's unique position in the vocabulary.

100

- For example, given a vocabulary of 5 words:
 - apple = [1, 0, 0, 0, 0]
 - banana = [0, 1, 0, 0, 0]
 - cherry = [0, 0, 1, 0, 0]
 - date = [0, 0, 0, 1, 0]
 - elderberry = [0, 0, 0, 0, 1]
- The length of the one-hot vector depends on the number of unique words in the vocabulary.

- One-hot encoding is a simple and intuitive representation that can be effective in certain models, especially smaller neural networks.
- It requires minimal computation and works well for small vocabularies or categorical features in simpler tasks.

- One-hot encoding does not capture any semantic relationships between words.
- The vectors are sparse, containing mostly zeros, which is inefficient for large vocabularies.
- Similar words (like `hotel` and `motel`) appear completely unrelated in this representation.

- **Cosine Similarity:**

$$\cos(\theta) = \frac{\mathbf{v}_1 \cdot \mathbf{v}_2}{\|\mathbf{v}_1\| \|\mathbf{v}_2\|}$$

Conclusion: Why Move Beyond One-Hot?

- While one-hot encoding is a simple and effective method for certain applications, it fails to capture word meanings or relationships.
- More advanced methods, such as word embeddings, address these limitations by representing words in a dense, meaningful vector space.

1 Introduction

② Word2Vec

Continuous Bag of Words (CBOW)

Skip-gram Model

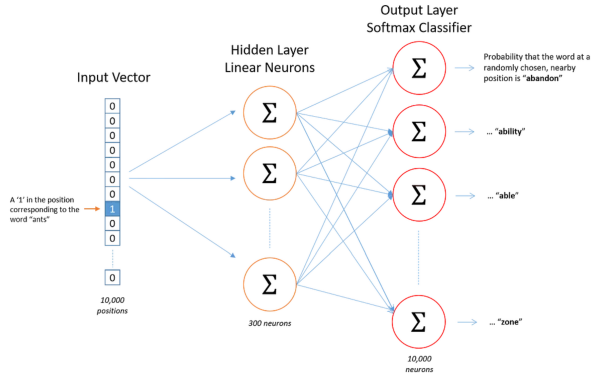
Word Embedding Visualization

Word Analogy

3 References

- To process text data, we need to represent words in a form that a machine can understand—numerical vectors.
- Word2Vec uses a neural network to learn word embeddings that capture semantic similarities.
- These embeddings allow words with similar meanings to be represented by vectors close to each other in a high-dimensional space.

- Word2Vec operates like a shallow neural network, with an input, hidden, and output layer.
- It takes in a target word and learns to predict either the surrounding context words or the target word from a set of context words.
- Through training, the network adjusts weights to create meaningful vector representations of words.



Word2Vec as a two layer neural network

Figure adapted from medium.com/@manansuri/a-dummys-guide-to-word2vec-456444f3c673

Expected Outcome of Word2Vec

- Word2Vec aims to create a vector space where words with similar meanings or contexts are located close to each other.
- **Expected Result:** Semantically related words—such as “king” and “queen” or “dog” and “puppy”—should have similar vector representations.
- This proximity allows for various NLP tasks, such as:
 - **Synonym detection:** Identifying words with similar meanings.
 - **Analogy tasks:** Solving analogies by vector arithmetic (e.g., “king” - “man” + “woman” “queen”).
 - **Clustering of concepts:** Grouping related concepts together in the embedding space.
- By representing words in this way, Word2Vec enables models to make use of semantic relationships between words.

100



Hypothetical features to understand word embeddings

Figure adapted from medium.com/@manansuri/a-dummys-guide-to-word2vec-456444f3c673

- The core idea is based on distributional semantics: "You shall know a word by the company it keeps."
- Word2Vec uses two main algorithms for learning word vectors:
 - Continuous Bag of Words (CBOW)
 - Skip-gram Model

CBOW: How It Works

- CBOW predicts the target word using the context (surrounding words) in a fixed window.
- For each word in the corpus, CBOW takes a set of context words and predicts the center word.
- Example: Given the context words {"the", "brown", "fox", "over"}, CBOW predicts the center word "jumps."
- CBOW tends to perform better on smaller datasets and is computationally more efficient.

① Introduction

② Word2Vec

Continuous Bag of Words (CBOW)

Skip-gram Model

Word Embedding Visualization

Word Analogy

③ References

Skip-gram: How It Works

- Skip-gram is the reverse of CBOW. It predicts the surrounding context words given a target word.
- For each word w_t , the model predicts the words in the window of size m around it (e.g., words $w_{t-2}, w_{t-1}, w_{t+1}, w_{t+2}$).
- Example: If the center word is “jumps,” Skip-gram predicts the context words “the,” “brown,” “fox,” and “over.”
- Skip-gram is better suited for larger datasets and can capture rare words more effectively.

Skip-gram Example

Window Size	Text	Skip-grams
2	[The wide road shimmered] in the hot sun.	wide, the wide, road wide, shimmered
	The [wide road shimmered in the] hot sun.	shimmered, wide shimmered, road shimmered, in shimmered, the
	The wide road shimmered in [the hot sun].	sun, the sun, hot
3	[The wide road shimmered in] the hot sun.	wide, the wide, road wide, shimmered wide, in
	[The wide road shimmered in the hot] sun.	shimmered, the shimmered, wide shimmered, road shimmered, in shimmered, the shimmered, hot
	The wide road shimmered [in the hot sun].	sun, in sun, the sun, hot

Different window sizes and samples drawn from context words and their target

Skip-gram: Objective Function

- The objective of Skip-gram is to maximize the likelihood of predicting context words w_o given a center word w_c .
- The probability of a context word w_o given a center word w_c is defined as:

$$P(w_o|w_c) = \frac{\exp(v_{w_o} \cdot v_{w_c})}{\sum_{w \in V} \exp(v_w \cdot v_{w_c})}$$

- v_{w_o} and v_{w_c} are the word vectors for the context and center words, respectively.

Skip-gram: Loss Function

- The goal is to minimize the negative log-likelihood over the entire training corpus:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} | w_t)$$

- Here, T is the total number of words, and m is the window size.
- Skip-gram adjusts the word vectors to maximize the probability of observing the context words around the center word.

Skip-gram: Gradient Calculation

- To update the word vectors during training, we calculate the gradient of the objective function.
- The gradient with respect to the word vector v_I is:

$$\frac{\partial \log P(w_o | w_I)}{\partial v_I} = u_o - \sum_x P(w_x | w_I) u_x$$

Skip-gram: Gradient Calculation

- The detailed steps for the gradient calculation are:

$$\begin{aligned}
 \frac{\partial \log P(w_o | w_I)}{\partial v_I} &= \frac{\partial}{\partial v_I} \log \frac{e^{u_o^T v_I}}{\sum_x e^{u_x^T v_I}} \\
 &= \frac{\partial}{\partial v_I} \left(\log e^{u_o^T v_I} - \log \sum_x e^{u_x^T v_I} \right) \\
 &= u_o - \frac{1}{\sum_x e^{u_x^T v_I}} \sum_x u_x e^{u_x^T v_I} \\
 &= u_o - \sum_x P(w_x | w_I) u_x
 \end{aligned}$$

- The update rule for v_{w_c} is:

$$v_{w_c} \leftarrow v_{w_c} + \eta \left(v_{w_o} - \sum_{w \in V} P(w | w_c) v_w \right)$$

- Here, η is the learning rate.

Skip-gram Example

- Consider the sentence: "The quick brown fox jumps over the lazy dog."
- If the center word is "jumps", the model predicts context words such as "quick", "brown", "fox", "over", and "the".
- The model iteratively adjusts word vectors to predict these context words, learning a meaningful representation for "jumps."

Skip-Gram Pseudocode

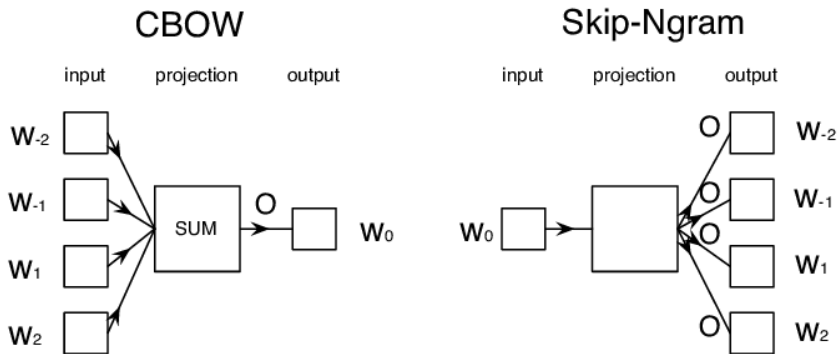
Algorithm 1 Skip-Gram Model

Require: Corpus D , window size w , embedding dimension d , learning rate α , number of epochs n
 Initialize word embeddings W and C randomly, where W maps words to embeddings and C maps context words to embeddings
for each epoch in 1 to n **do**
 for each sentence S in D **do**
 for each word w_t in S **do**
 Extract context words within window size w around w_t
 for each context word c of w_t **do**
 Compute dot product score = $W(w_t) \cdot C(c)$
 Compute probability $P(c|w_t)$ using softmax:

$$P(c|w_t) = \frac{\exp(\text{score})}{\sum_{c' \in \text{vocab}} \exp(W(w_t) \cdot C(c'))}$$

 Calculate loss $L = -\log P(c|w_t)$
 Update $W(w_t)$ and $C(c)$ using gradient descent with learning rate η
 end for
 end for
end for
return Word embeddings W

CBOW vs. Skip-gram



Difference between CBOW and Skip-gram

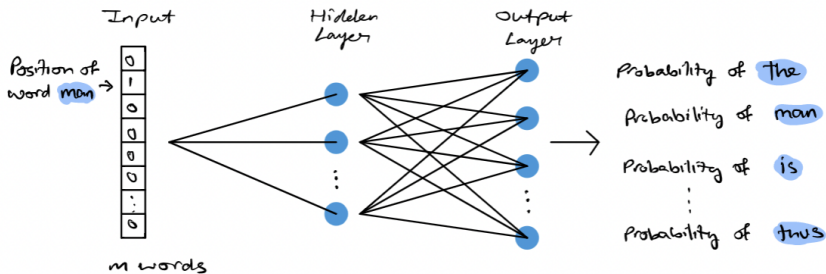
CBOW vs. Skip-gram

- CBOW and Skip-gram are the two primary architectures for Word2Vec.
- **CBOW** predicts a target word given its surrounding context, making it efficient and effective for smaller datasets.
- **Skip-gram**, on the other hand, predicts the surrounding words for a given target word. It is well-suited for larger datasets and can handle rare words more effectively.
- In essence, the Skip-gram model captures more detailed word relationships and is robust in large vocabularies.

Skip-gram as a Neural Network

- The Skip-gram model functions as a two-layer neural network.
- The input layer consists of a one-hot encoded target word vector, while the hidden layer is a dense embedding layer that learns the word representation.
- The output layer uses softmax to calculate the probability distribution over all words in the vocabulary, given the context.
- The Skip-gram model iteratively adjusts weights to maximize the likelihood of predicting the correct context words, ultimately learning meaningful word embeddings.

Skip-gram as a Neural Network



Skip-gram Model as Neural Network

Figure adapted from towardsdatascience.com/word2vec-explained-49c52b4ccb71

Why Do We Need Negative Sampling?

- The softmax function normalizes over all words in the vocabulary V , which can be very large (millions of words).
- This makes the calculation of $\sum_{w \in V} \exp(v_w \cdot v_{w_c})$ expensive, as it requires summing over all words in the vocabulary.
- **Solution:** Use **Negative Sampling** to only update a few randomly chosen "negative" words instead of the entire vocabulary.

Skip-gram: Negative Sampling

- In **negative sampling**, we sample a few "negative" words that do not appear in the context of the target word.
- For each positive pair (center word and context word), we sample k negative words that are not in the context.
- Instead of maximizing the probability of all words in the vocabulary, we only maximize the probability of the context words and minimize the probability of the sampled negative words.
- Example: If the center word is "cat", and the context word is "cute", we sample negative words like "computer", "sky", and "table" to minimize their probability in this context.

Skip-gram: Negative Sampling

Skipgram

shalt	not	make	a	machine
input		output		
make		shalt		
make		not		
make		a		
make		machine		

Negative Sampling

input word	output word	target
make	shalt	1
make	aaron	0
make	taco	0

Negative sampling

Figure adapted from alammar.github.io/illustrated-word2vec/

Why Skip-gram?

- Skip-gram is preferred for large datasets because it handles rare words more effectively than CBOW.
- It captures detailed information about the surrounding words, leading to better word representations in the vector space.
- Word2Vec embeddings from Skip-gram have been widely adopted in various NLP tasks such as machine translation, sentiment analysis, and document classification.

① Introduction

② Word2Vec

Continuous Bag of Words (CBOW)

Skip-gram Model

Word Embedding Visualization

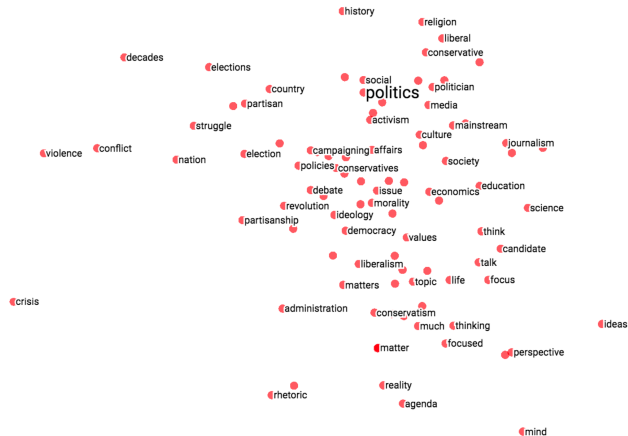
Word Analogy

③ References

Visualizing Words in 2D

- After training the model, words are mapped into a high-dimensional vector space.
- Using techniques like PCA or t-SNE, these vectors can be reduced to 2D for visualization, where similar words appear closer together.

Visualizing Words in 2D



Words represented in a 2D space after dimensionality reduction

① Introduction

② Word2Vec

Continuous Bag of Words (CBOW)

Skip-gram Model

Word Embedding Visualization

Word Analogy

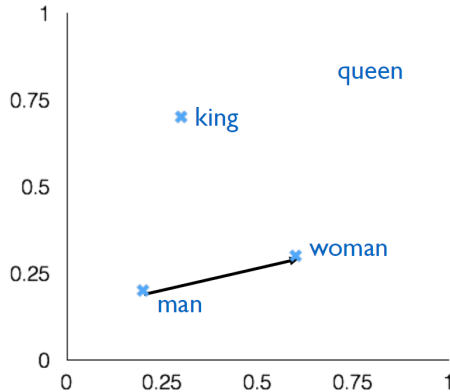
③ References

Word Analogy: Vector Arithmetic in Word2Vec

- Word2Vec embeddings can solve analogy tasks by performing vector arithmetic.
- The analogy task takes the form:

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

- The analogy is solved by finding the word vector closest to $\mathbf{v}_{king} - \mathbf{v}_{man} + \mathbf{v}_{woman}$.



Word analogy example

- $$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}$$
- This finds the word x_i whose vector is closest to the result of the vector arithmetic.

$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}$$

- This finds the word x_i whose vector is closest to the result of the vector arithmetic.

① Introduction

② Word2Vec

③ References

- [1] M. Soleymani, “Machine learning.” Sharif University of Technology.
- [2] M. Soleymani, “Modern information retrieval.” Sharif University of Technology.
- [3] C. Manning, “Natural language processing with deep learning.” Stanford.
- [4] E. Asgari, “Natural language processing.” Sharif University of Technology.