



Encoder-only (e.g., BERT): bidirectional contextual embeddings



Decoder-only (e.g., GPT-x): unidirectional contextual embeddings, generate one token at a time



Encoder-decoder (e.g., T5): encode input, decode output